

Analysis Report: iCEVOS Process Scheduler

Componentes:
Paulo Vinicius Alves Melo

Teresina-PI
09/2025

1. Design Justification

Why is a manual linked list efficient for the scheduler?

1. $O(1)$ insertion at tail – perfect for FIFO queues
2. $O(1)$ removal from head – matches the “next to run” semantics
3. No random access needed – we never need the n -th process
4. Dynamic size – no waste of pre-allocated memory
5. Simple code – easy to debug and prove correctness

Using an array or vector would require $O(n)$ shifts or amortized resizing, making linked lists the natural, efficient choice for this problem.

2. Big-O Complexity Analysis

| Operation | Complexity | Explanation |
|----------------------------|--------------------------|--------------------------------|
| insert_end(process) | $O(1)$ | Tail pointer update |
| remove_start() | $O(1)$ | Head pointer update |
| run_cpu_cycles() | $O(1)$ | Constant work per cycle |
| unblock() | $O(1)$ | Single remove + append |
| load_process() | $O(n)$ | One pass over file |

All scheduler-critical operations are constant time, ensuring scalability even with thousands of processes.

3. Anti-Starvation Analysis

Rule: After 5 consecutive HIGH executions, force one MEDIUM (or LOW if MEDIUM is empty).

Fairness guarantee:

- Every 5 cycles lower priorities get a chance
- Bounded waiting time = $5 \times$ shortest HIGH job
- No process waits indefinitely

Risk without the rule:

LOW/MEDIUM processes could starve indefinitely if HIGH queue never empties (e.g., constant arrival of HIGH jobs).

The rule introduces predictable fairness with minimal impact on HIGH throughput.

4. Disc Blocking Life-Cycle

1. Arrival – process enters its priority queue
2. First selection – scheduler detects `resource_needed == "DISC"` and `was_blocked == false`
3. Block – process is removed from priority queue and appended to blocked list
4. Wait – process stays in blocked while others execute
5. Unblock – at the start of the next cycle, the oldest blocked process is moved back to the tail of its original priority queue
6. Re-execution – process is now eligible to run; if it still needs DISC, it will not be blocked again (flag already set)
7. Completion – cycles reach 0 → process terminates

5. Performance Bottleneck & Theoretical Improvement

Main bottleneck: Global starvation counter and strict FIFO queues – a burst of HIGH jobs can still delay MEDIUM/LOW for up to 5 cycles, and aging is not supported.

Theoretical improvement – Aging with Multi-Level Feedback Queue (MLFQ):

- Add priority boost bit: increment a wait-counter each cycle a process is skipped
- When counter > threshold, promote the process one level (e.g., LOW → MEDIUM)
- Reset counter after promotion
- Complexity still $O(1)$ per operation, but starvation bound becomes 1 cycle instead of 5
- Throughput of HIGH jobs remains high, latency of LOW jobs drops significantly

Conclusion:

The manual linked-list scheduler is simple, fair, and efficient for the proposed scenario.

The anti-starvation rule and “disc” blocking mechanism ensure correctness and resource safety.

Aging-enhanced MLFQ is the next evolutionary step for production-grade systems.